

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import OrdinalEncoder
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error as MSE
from sklearn.metrics import mean_absolute_error as mae
from sklearn.metrics import roc_auc_score

path='/Users/herlihpj/Desktop/Data Analytics/D209 - Data Mining/Task 2/'

#Data Preparation
#Reads CSV to data frame, sets case order to index
med_dirty= pd.read_csv(path+'medical_clean.csv',
                      index_col=0)

#Check for Null
print('Summary of Null: ')
print(med_dirty.isna().sum())

#Check for duplicated data
duplicates=med_dirty.duplicated()
print('Duplicates: ', duplicates.sum())

# Dropping columns not relevant to the analysis
med_mine = med_dirty.drop(columns= ["Customer_id", "Interaction", 'TimeZone', "UID", "City", "County", "Zip", "Lat",
                                   "Lng", 'Job', 'Item1', 'Item2', 'Item3', 'Item4', 'Item5', 'Item6', 'Item7', 'Item8'])

#Explore the data/statistics
print(med_mine.head())
print(med_mine.describe())
print(med_mine.info())

#Various Scatterplots to Visualize data
sns.scatterplot(data=med_mine, x="Initial_days", y="TotalCharge", hue='Initial_admin')
plt.show()
sns.scatterplot(data=med_mine, x="Initial_days", y="VitD_levels", hue='ReAdmis')
plt.show()

#Ordinal Encoding to convert to numeric 0:No, 1:Yes; other variable alphabetically starting with 0
oe_dict={}
#List of columns to convert to numerical
```

```

convert_cols=['State','Area','Services', 'Marital', 'Gender','Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke',
              'Complication_risk','Overweight', 'Arthritis', 'Diabetes', 'Hyperlipidemia','BackPain', 'Anxiety',
              'Allergic_rhinitis', 'Reflux_esophagitis', 'Asthma', 'Services','ReAdmis']#
for col_name in convert_cols:
    #print(col_name+' pre: '+str(med_mine[col_name].unique()))
    #Creates column ordinal encoder
    oe_dict[col_name]=OrdinalEncoder()
    col=med_mine[col_name]
    #select non-null values of col
    col_not_null=col[col.notnull()]
    reshaped_vals=col_not_null.values.reshape(-1,1)
    encoded_vals=oe_dict[col_name].fit_transform(reshaped_vals)
    med_mine.loc[col.notnull(), col_name]=np.squeeze(encoded_vals)
    #print(col_name+' post: '+str(med_mine[col_name].unique()))

#Visual EDA - wouldnt plot with color y?
#_= pd.plotting.scatter_matrix(med_mine, c = 'green', figsize = [8, 8],s=150, marker = 'D')
#returns a series of plots and histograms
#Slow runtime too many graphs with this data

med_mine.to_csv(path+'Prepared_data.csv')
print('Prepared Data has been exported to CSV')

print('===== \n Data has been prepared      \n===== ')

#Function which takes y_test values, prediction values, and a range to compare against
#Function scores predictions based on accuracy
def score_within(test, predictions, buffer):
    score=0
    for x, pred in enumerate(predictions):
        test_min=test.iloc[x]-buffer
        test_max=test.iloc[x]+buffer
        #print('Correct: ',y_test.iloc[x], ' Predicted: ', pred)
        if pred>test_min and pred<test_max:
            score=score+1
    return score

### RANDOM FOREST - Hyperparameter Tuned REGRESSION with Grid Search Cross Validation ###
print('RANDOM FORESTS:')

#Set the target variable
target='Initial_days'

#Drop additional columns not necessary to the analysis
med_mine=med_mine.drop(columns = ['Additional_charges'])# 'TotalCharge', 'ReAdmis',

```

```
#Train/test Split
X=med_mine.drop(target, axis=1)#.values #can drop readmis too doesnt change anything
y=med_mine[target]
#Check to make sure target has same numer of rows and just one value
print(X.shape) #(10000, 27)
print(y.shape) #(10000,)

# Set seed for reproducibility
SEED = 1
# Split dataset into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=SEED)

#Export split data
# Convert X_train and y_train into DataFrames
X_train_df = pd.DataFrame(X_train)
y_train_df = pd.DataFrame(y_train)
X_test_df = pd.DataFrame(X_test)
y_test_df = pd.DataFrame(y_test)
# Merge X_train_df and y_train_df into a single DataFrame
train = X_train_df.join(y_train_df)
test = X_test_df.join(y_test_df)
#Export to CSV's
train.to_csv(path+'Training_data.csv')
test.to_csv(path+'Testing_data.csv')
print('Split data has been exported')

# Instantiate a random forests regressor 'rf'
rf = RandomForestRegressor(random_state= SEED)
# Inspect rf' s hyperparameters
print(rf.get_params())

params_rf = {'n_estimators': [100, 200, 300],
             'max_depth': [4, 6, 8],
             'min_samples_leaf': [0.08, 0.1, 0.15],
             'max_features': ['log2', 'sqrt']}

# Instantiate 'grid_rf'
grid_rf = GridSearchCV(estimator=rf, param_grid=params_rf, cv=3, scoring='neg_mean_squared_error', verbose=1, n_jobs=-1)

#Search for the Best Hyper Parameters
# Fit 'grid_rf' to the training set
grid_rf.fit(X_train, y_train)
#output shows messages to grid fitting and the obtained optimal model
```

```

# Extract the best hyperparameters from 'grid_rf'
best_hyperparams = grid_rf.best_params_
print('Best hyperparameters:\n', best_hyperparams)

# Extract the best model from 'grid_rf'
best_model = grid_rf.best_estimator_
# Predict the test set labels
y_pred = best_model.predict(X_test)

# Feature Importance
# Create a pd.Series of features importances
importances_rf = pd.Series(data=best_model.feature_importances_, index=X.columns)
# Sort importances_rf
sorted_importances_rf = importances_rf.sort_values()
# Make a horizontal bar plot
sorted_importances_rf.plot(kind='barh', color='lightgreen'); plt.show()
#Plots a horizontal bar graph with each features importance

#Can also store and sort the importances first then loop through
#print('Top Features Causing '+target+':')
#for i, item in enumerate(sorted_importances_rf):
#    if (item>.05):
#        print("{0:s}: {1:.2f}".format(X.columns[i], item))

days_within=14
print('Score: ', score_within(y_test, y_pred, days_within), ' Of: ', len(y_pred))

# Evaluate the train & test set MAE
##
print('Train vs Test MAEs:')
# Create vectors of predictions
train_predictions = best_model.predict(X_train)
test_predictions = best_model.predict(X_test)
# Train/Test Errors
train_error = mae(y_true=y_train, y_pred=train_predictions)
test_error = mae(y_true=y_test, y_pred=test_predictions)
# Print the accuracy for seen and unseen data
print("Model error on seen data: {0:.2f}.".format(train_error))
print("Model error on unseen data: {0:.2f}.".format(test_error))
##

# Evaluate the train & test set MSE and RMSE
mse_test=MSE(y_test, y_pred)
rmse_test = mse_test**(1/2)
# Print the test set MSE

```

```
print('Test set MSE: {:.2f}'.format(mse_test))  
# Print the test set RMSE  
print('Test set RMSE: {:.2f}'.format(rmse_test))  
#Test set  
print("Test Set MAE: {0:.2f}".format(mae(y_true=y_test, y_pred=y_pred)))
```

Summary of Null:

Customer_id	0
Interaction	0
UID	0
City	0
State	0
County	0
Zip	0
Lat	0
Lng	0
Population	0
Area	0
TimeZone	0
Job	0
Children	0
Age	0
Income	0
Marital	0
Gender	0
ReAdmis	0
VitD_levels	0
Doc_visits	0
Full_meals_eaten	0
vitD_supp	0
Soft_drink	0
Initial_admin	0
HighBlood	0
Stroke	0
Complication_risk	0
Overweight	0
Arthritis	0
Diabetes	0
Hyperlipidemia	0
BackPain	0
Anxiety	0
Allergic_rhinitis	0
Reflux_esophagitis	0
Asthma	0
Services	0
Initial_days	0
TotalCharge	0
Additional_charges	0
Item1	0
Item2	0
Item3	0

```
Item4      0
Item5      0
Item6      0
Item7      0
Item8      0
```

```
dtype: int64
```

```
Duplicates: ; 0
```

	State	Population	Area	Children	Age	Income	Marital	\
CaseOrder								
1	AL	2951	Suburban	1	53	86575.93	Divorced	
2	FL	11303	Urban	3	51	46805.99	Married	
3	SD	17125	Suburban	3	53	14370.14	Widowed	
4	MN	2162	Suburban	0	78	39741.49	Married	
5	VA	5287	Rural	1	22	1209.56	Widowed	

	Gender	ReAdmis	VitD_levels	...	Hyperlipidemia	BackPain	\
CaseOrder				...			
1	Male	No	19.141466	...	No	Yes	
2	Female	No	18.940352	...	No	No	
3	Female	No	18.057507	...	No	No	
4	Male	No	16.576858	...	No	No	
5	Female	No	17.439069	...	Yes	No	

	Anxiety	Allergic_rhinitis	Reflux_esophagitis	Asthma	Services	\
CaseOrder						
1	Yes		Yes	No	Yes	Blood Work
2	No		No	Yes	No	Intravenous
3	No		No	No	No	Blood Work
4	No		No	Yes	Yes	Blood Work
5	No		Yes	No	No	CT Scan

	Initial_days	TotalCharge	Additional_charges
CaseOrder			
1	10.585770	3726.702860	17939.403420
2	15.129562	4193.190458	17612.998120
3	4.772177	2434.234222	17505.192460
4	1.714879	2127.830423	12993.437350
5	1.254807	2113.073274	3716.525786

```
[5 rows x 31 columns]
```

	Population	Children	Age	Income	VitD_levels	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	9965.253800	2.097200	53.511700	40490.495160	17.964262	
std	14824.758614	2.163659	20.638538	28521.153293	2.017231	
min	0.000000	0.000000	18.000000	154.080000	9.806483	

25%	694.750000	0.000000	36.000000	19598.775000	16.626439
50%	2769.000000	1.000000	53.000000	33768.420000	17.951122
75%	13945.000000	3.000000	71.000000	54296.402500	19.347963
max	122814.000000	10.000000	89.000000	207249.100000	26.394449

	Doc_visits	Full_meals_eaten	vitD_supp	Initial_days	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	5.012200	1.001400	0.398900	34.455299	
std	1.045734	1.008117	0.628505	26.309341	
min	1.000000	0.000000	0.000000	1.001981	
25%	4.000000	0.000000	0.000000	7.896215	
50%	5.000000	1.000000	0.000000	35.836244	
75%	6.000000	2.000000	1.000000	61.161020	
max	9.000000	7.000000	5.000000	71.981490	

	TotalCharge	Additional_charges
count	10000.000000	10000.000000
mean	5312.172769	12934.528587
std	2180.393838	6542.601544
min	1938.312067	3125.703000
25%	3179.374015	7986.487755
50%	5213.952000	11573.977735
75%	7459.699750	15626.490000
max	9180.728000	30566.070000

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 10000 entries, 1 to 10000
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	State	10000 non-null	object
1	Population	10000 non-null	int64
2	Area	10000 non-null	object
3	Children	10000 non-null	int64
4	Age	10000 non-null	int64
5	Income	10000 non-null	float64
6	Marital	10000 non-null	object
7	Gender	10000 non-null	object
8	ReAdmis	10000 non-null	object
9	VitD_levels	10000 non-null	float64
10	Doc_visits	10000 non-null	int64
11	Full_meals_eaten	10000 non-null	int64
12	vitD_supp	10000 non-null	int64
13	Soft_drink	10000 non-null	object
14	Initial_admin	10000 non-null	object
15	HighBlood	10000 non-null	object


```

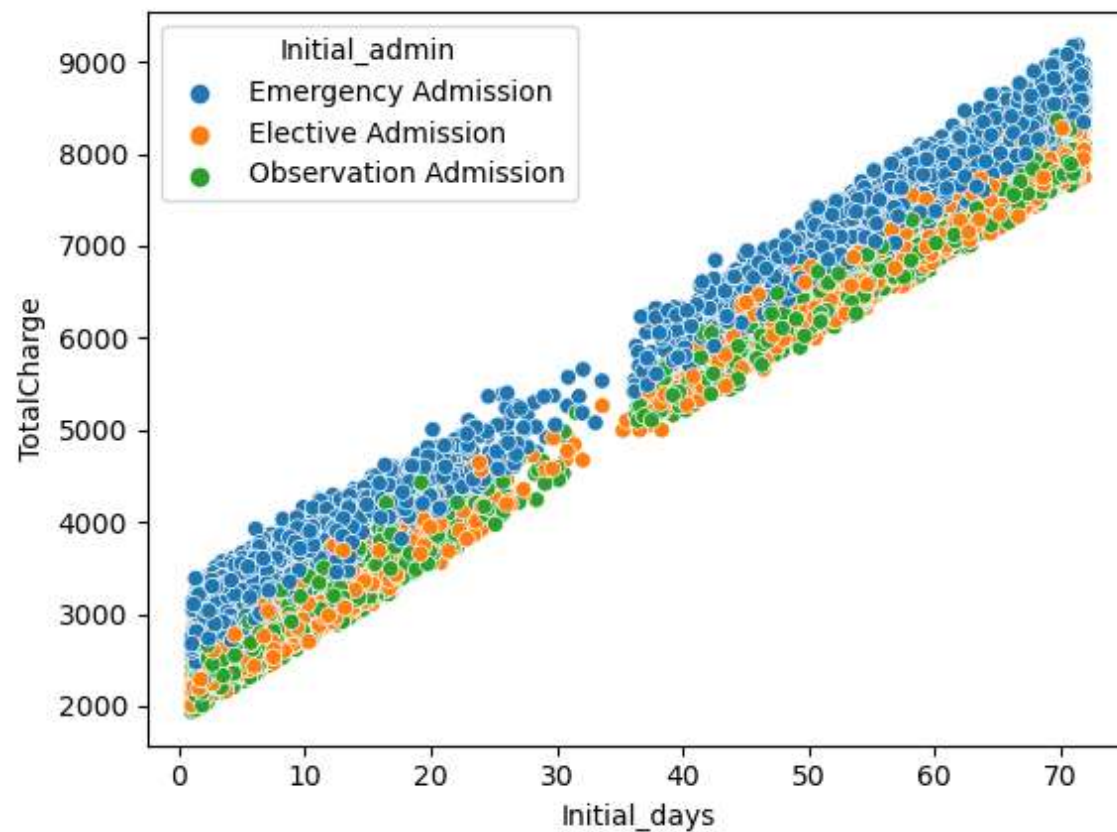
16 Stroke                10000 non-null object
17 Complication_risk     10000 non-null object
18 Overweight            10000 non-null object
19 Arthritis             10000 non-null object
20 Diabetes              10000 non-null object
21 Hyperlipidemia        10000 non-null object
22 BackPain              10000 non-null object
23 Anxiety               10000 non-null object
24 Allergic_rhinitis     10000 non-null object
25 Reflux_esophagitis    10000 non-null object
26 Asthma                10000 non-null object
27 Services              10000 non-null object
28 Initial_days          10000 non-null float64
29 TotalCharge           10000 non-null float64
30 Additional_charges    10000 non-null float64


```

dtypes: float64(5), int64(6), object(20)

memory usage: 2.4+ MB

None





Prepared Data has been exported to CSV

=====

Data has been prepared

=====

RANDOM FORESTS:

(10000, 29)

(10000,)


Split data has been exported

```
{'bootstrap': True, 'ccp_alpha': 0.0, 'criterion': 'squared_error', 'max_depth': None, 'max_features': 'auto', 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'n_jobs': None, 'oob_score': False, 'random_state': 1, 'verbose': 0, 'warm_start': False}
```

Fitting 3 folds for each of 54 candidates, totalling 162 fits

Best hyperparameters:

```
{'max_depth': 6, 'max_features': 'sqrt', 'min_samples_leaf': 0.08, 'n_estimators': 300}
```



Score: 1955 Of: 3000

Train vs Test MAEs:

Model error on seen data: 11.31.

Model error on unseen data: 11.39.

Test set MSE: 155.91

Test set RMSE: 12.49

Test Set MAE: 11.39

In []: